

assignment-3-smartinternz

September 20, 2023

1 1. IMPORT THE LIBRARIES

```
[5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

2 2. IMPORT THE DATASET

```
[6]: df=pd.read_csv("Titanic-Dataset.csv")
```

```
[7]: df
```

```
[7]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	
..	
886	887	0	2	
887	888	1	1	
888	889	0	3	
889	890	1	1	
890	891	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	
..	

886		Montvila, Rev. Juozas	male	27.0	0
887		Graham, Miss. Margaret Edith	female	19.0	0
888		Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1
889		Behr, Mr. Karl Howell	male	26.0	0
890		Dooley, Mr. Patrick	male	32.0	0

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
..
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	NaN	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

```
[8]: df.head()
```

```
[8]: PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
3             4         1         1
4             5         0         3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[9]: df.tail()
```

```
[9]:
```

	PassengerId	Survived	Pclass	Name \
886	887	0	2	Montvila, Rev. Juozas
887	888	1	1	Graham, Miss. Margaret Edith
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"
889	890	1	1	Behr, Mr. Karl Howell
890	891	0	3	Dooley, Mr. Patrick

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
886	male	27.0	0	0	211536	13.00	NaN	S
887	female	19.0	0	0	112053	30.00	B42	S
888	female	NaN	1	2	W./C. 6607	23.45	NaN	S
889	male	26.0	0	0	111369	30.00	C148	C
890	male	32.0	0	0	370376	7.75	NaN	Q

```
[10]: df.shape
```

```
[10]: (891, 12)
```

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age             714 non-null    float64
6   SibSp           891 non-null    int64
7   Parch           891 non-null    int64
8   Ticket          891 non-null    object
9   Fare            891 non-null    float64
10  Cabin           204 non-null    object
11  Embarked        889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[12]: df.describe()
```

```
[12]:
```

	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	714.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008
std	257.353842	0.486592	0.836071	14.526497	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000

25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
[13]: corr=df.corr()
corr
```

<ipython-input-13-7d5195e2bf4d>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr=df.corr()
```

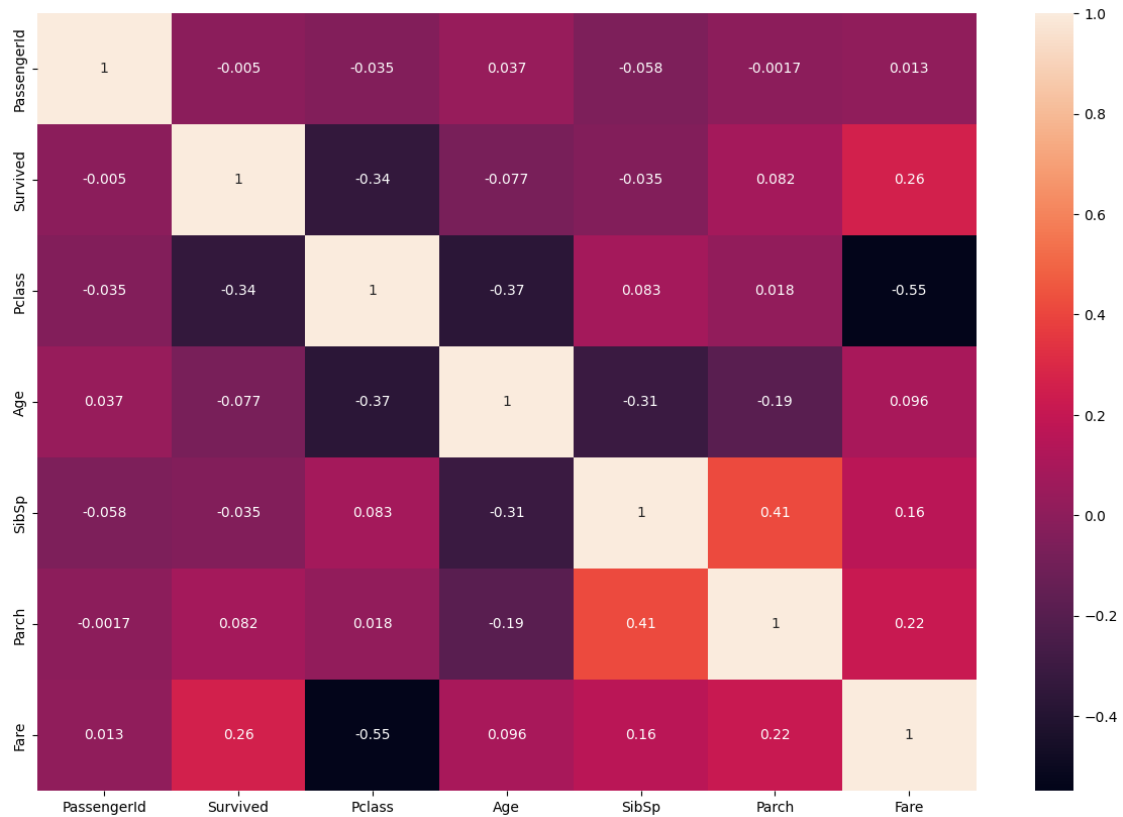
```
[13]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	\
PassengerId	1.000000	-0.005007	-0.035144	0.036847	-0.057527	-0.001652	
Survived	-0.005007	1.000000	-0.338481	-0.077221	-0.035322	0.081629	
Pclass	-0.035144	-0.338481	1.000000	-0.369226	0.083081	0.018443	
Age	0.036847	-0.077221	-0.369226	1.000000	-0.308247	-0.189119	
SibSp	-0.057527	-0.035322	0.083081	-0.308247	1.000000	0.414838	
Parch	-0.001652	0.081629	0.018443	-0.189119	0.414838	1.000000	
Fare	0.012658	0.257307	-0.549500	0.096067	0.159651	0.216225	

	Fare
PassengerId	0.012658
Survived	0.257307
Pclass	-0.549500
Age	0.096067
SibSp	0.159651
Parch	0.216225
Fare	1.000000

```
[14]: plt.subplots(figsize=(15,10))
sns.heatmap(corr,annot=True)
```

```
[14]: <Axes: >
```



```
[15]: df.Survived.value_counts()
```

```
[15]: 0    549
      1    342
      Name: Survived, dtype: int64
```

```
[16]: df.Sex.value_counts()
```

```
[16]: male    577
      female  314
      Name: Sex, dtype: int64
```

```
[17]: df.Embarked.value_counts()
```

```
[17]: S    644
      C    168
      Q     77
      Name: Embarked, dtype: int64
```

#3. CHECK FOR NULL VALUES

```
[18]: df.isnull().any()
```

```
[18]: PassengerId    False
      Survived      False
      Pclass        False
      Name          False
      Sex           False
      Age           True
      SibSp         False
      Parch         False
      Ticket        False
      Fare          False
      Cabin         True
      Embarked      True
      dtype: bool
```

```
[19]: df.isnull().sum()
```

```
[19]: PassengerId      0
      Survived        0
      Pclass          0
      Name            0
      Sex             0
      Age            177
      SibSp           0
      Parch           0
      Ticket          0
      Fare            0
      Cabin          687
      Embarked        2
      dtype: int64
```

Fill null values in the 'Age' column with the mean age

```
[20]: mean_age = df['Age'].mean()
      df['Age'].fillna(mean_age, inplace=True)
```

Fill null values in the 'Embarked' column with the most common value

```
[21]: most_common_embarked = df['Embarked'].mode()[0]
      df['Embarked'].fillna(most_common_embarked, inplace=True)
```

```
[22]: df.drop(['Cabin'],axis=1, inplace=True)
```

```
[23]: df.drop(['Ticket'],axis=1, inplace=True)
```


```
[24]: df.drop(['Name'],axis=1,inplace=True)
```

```
[25]: print(df.isnull().sum())
```


```
PassengerId    0
Survived        0
Pclass          0
Sex             0
Age            0
SibSp           0
Parch           0
Fare            0
Embarked        0
dtype: int64
```

#4. Data Visualization

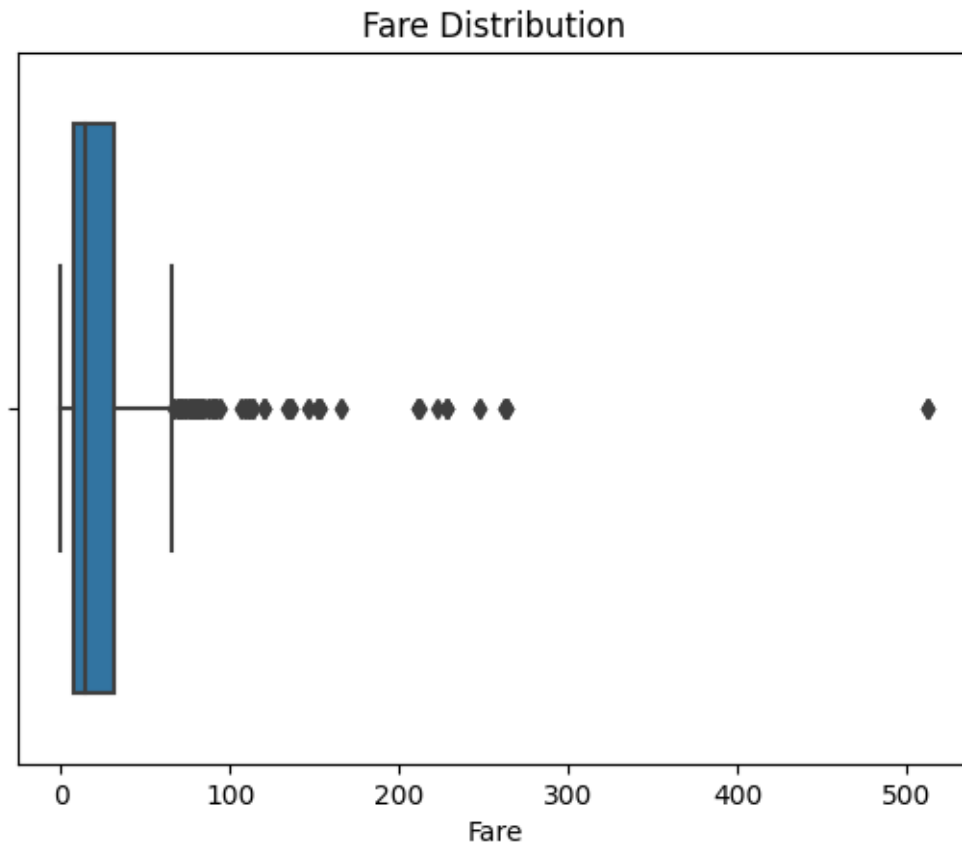
```
[26]: # Visualize the distribution of the 'Survived' column (0 = Not Survived, 1 = Survived)
sns.countplot(data=df, x='Survived')
plt.title('Survival Count')
plt.xlabel('Survived')
plt.ylabel('Count')
plt.show()
```



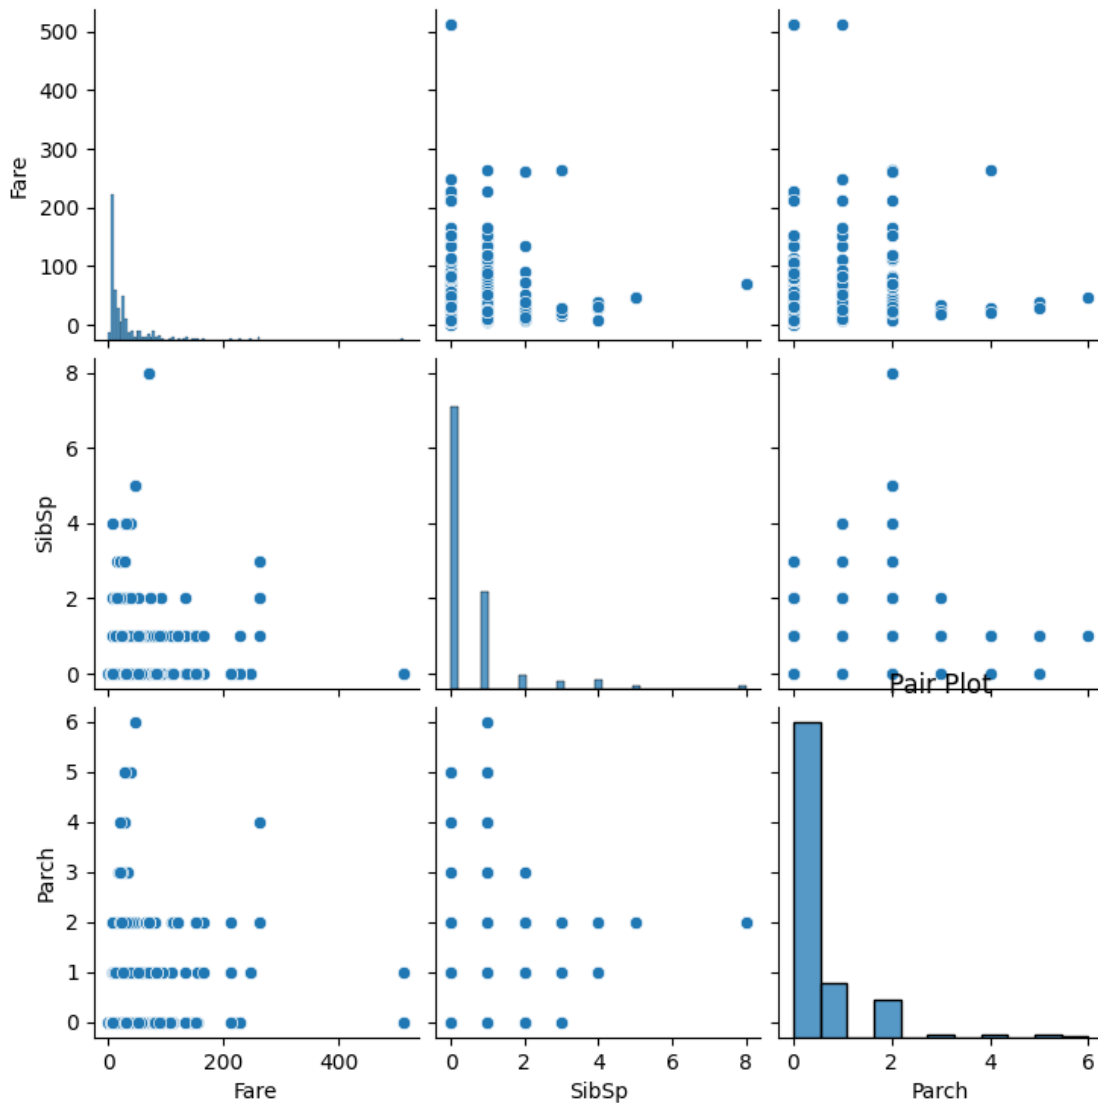
```
[27]: #Visualize the distribution of the 'Age' column
sns.histplot(data=df, x='Age', bins=20, kde=True)
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```



```
[28]: #Visualize the distribution of the 'Fare' column and detect outliers we will
      ↪ handle outliers in the next step
sns.boxplot(data=df, x='Fare')
plt.title('Fare Distribution')
plt.xlabel('Fare')
plt.show()
```


```
[29]: #Pair plot for selected numerical columns  
sns.pairplot(data=df[['Fare', 'SibSp', 'Parch']])  
plt.title('Pair Plot')  
plt.show()
```



```
[30]: corr_matrix = df.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```

<ipython-input-30-8dcbd071fff3>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr_matrix = df.corr()
```



3 5. Detect and Handle Outliers

```
[31]: z_scores = np.abs(stats.zscore(df['Age']))
max_threshold=3
outliers = df['Age'][z_scores > max_threshold]

# Print and visualize the outliers
print("Outliers detected using Z-Score:")
print(outliers)
```

Outliers detected using Z-Score:

```
96      71.0
116     70.5
493     71.0
630     80.0
672     70.0
745     70.0
```

```
851      74.0
Name: Age, dtype: float64
```

```
[32]: z_scores = np.abs(stats.zscore(df['Fare']))
      max_threshold=3
      outliers = df['Fare'][z_scores > max_threshold]

      # Print and visualize the outliers
      print("Outliers detected using Z-Score:")
      print(outliers)
```

```
Outliers detected using Z-Score:
```

```
27      263.0000
88      263.0000
118     247.5208
258     512.3292
299     247.5208
311     262.3750
341     263.0000
377     211.5000
380     227.5250
438     263.0000
527     221.7792
557     227.5250
679     512.3292
689     211.3375
700     227.5250
716     227.5250
730     211.3375
737     512.3292
742     262.3750
779     211.3375
```

```
Name: Fare, dtype: float64
```

```
[33]: column_name = 'Fare'

      # Calculate the first quartile (Q1) and third quartile (Q3)
      Q1 = df[column_name].quantile(0.25)
      Q3 = df[column_name].quantile(0.75)

      # Calculate the IQR
      IQR = Q3 - Q1

      # Define the lower and upper bounds for outliers
      lower_bound = Q1 - 1.5 * IQR
      upper_bound = Q3 + 1.5 * IQR
```

```
# Filter rows with values outside the IQR bounds
df_cleaned = df[(df[column_name] > lower_bound) & (df[column_name]
↳upper_bound)]

# Display the original and cleaned DataFrame sizes
print(f"Original DataFrame size: {df.shape}")
print(f"Cleaned DataFrame size: {df_cleaned.shape}")
df_cleaned
```

Original DataFrame size: (891, 9)

Cleaned DataFrame size: (775, 9)

```
[33]:
```


	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	\
0	1	0	3	male	22.000000	1	0	7.2500	
2	3	1	3	female	26.000000	0	0	7.9250	
3	4	1	1	female	35.000000	1	0	53.1000	
4	5	0	3	male	35.000000	0	0	8.0500	
5	6	0	3	male	29.699118	0	0	8.4583	
..	
886	887	0	2	male	27.000000	0	0	13.0000	
887	888	1	1	female	19.000000	0	0	30.0000	
888	889	0	3	female	29.699118	1	2	23.4500	
889	890	1	1	male	26.000000	0	0	30.0000	
890	891	0	3	male	32.000000	0	0	7.7500	

```
Embarked
0      S
2      S
3      S
4      S
5      Q
..     ...
886    S
887    S
888    S
889    C
890    Q
```

[775 rows x 9 columns]

```
[34]: sns.boxplot(df_cleaned)
```

```
[34]: <Axes: >
```



```
[35]: df=df_cleaned
```

```
[36]: x=df.drop('Survived', axis=1)
      y=df['Survived']
```

```
[37]: x.head()
```

```
[37]:   PassengerId  Pclass   Sex    Age  SibSp  Parch   Fare Embarked
0            1         3  male  22.000000     1     0   7.2500         S
2            3         3 female  26.000000     0     0   7.9250         S
3            4         1 female  35.000000     1     0  53.1000         S
4            5         3  male  35.000000     0     0   8.0500         S
5            6         3  male  29.699118     0     0   8.4583         Q
```

```
[38]: y.head()
```

```
[38]: 0    0
      2    1
      3    1
      4    0
      5    0
      Name: Survived, dtype: int64
```

#7. Perform Encoding

```
[39]: en = LabelEncoder()  
x['Sex'] = en.fit_transform(x['Sex'])
```

```
[40]: x.head()
```

```
[40]:
```

	PassengerId	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	1	3	1	22.000000	1	0	7.2500	S
2	3	3	0	26.000000	0	0	7.9250	S
3	4	1	0	35.000000	1	0	53.1000	S
4	5	3	1	35.000000	0	0	8.0500	S
5	6	3	1	29.699118	0	0	8.4583	Q

```
[41]: x = pd.get_dummies(x,columns=['Embarked'])
```

```
[42]: x.head()
```

```
[42]:
```

	PassengerId	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked_C	\
0	1	3	1	22.000000	1	0	7.2500	0	
2	3	3	0	26.000000	0	0	7.9250	0	
3	4	1	0	35.000000	1	0	53.1000	0	
4	5	3	1	35.000000	0	0	8.0500	0	
5	6	3	1	29.699118	0	0	8.4583	0	

	Embarked_Q	Embarked_S
0	0	1
2	0	1
3	0	1
4	0	1
5	1	0

#8. Feature Scaling

```
[43]: scale = StandardScaler()  
x[['Age', 'Fare']] = scale.fit_transform(x[['Age', 'Fare']])
```

```
[44]: x.head()
```

```
[44]:
```

	PassengerId	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked_C	\
0	1	3	1	-0.556219	1	0	-0.779117	0	
2	3	3	0	-0.243027	0	0	-0.729373	0	
3	4	1	0	0.461654	1	0	2.599828	0	
4	5	3	1	0.461654	0	0	-0.720161	0	
5	6	3	1	0.046606	0	0	-0.690071	0	

	Embarked_Q	Embarked_S
0	0	1

2	0	1
3	0	1
4	0	1
5	1	0

#9. Splitting the data into Train and Test

```
[45]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,  
↳ random_state=42)
```

```
[46]: print(x_train.shape)  
print(x_test.shape)  
print(y_train.shape)  
print(y_test.shape)
```

(620, 10)

(155, 10)

(620,)

(155,)